

HAR Project

Shawnzoom

January 26, 2016

```
##
## Attaching package: 'dplyr'
##
## The following object is masked from 'package:stats':
##
##     filter
##
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

Using Sensor Data to Classify how well an Exercise Activity is Performed

Synopsis

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, the goal was to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants, to classify how well an activity was performed. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. The “raw” data had 19622 observations and 160 features. After data exploration and cleansing, the data was split into training and test set, and the feature count was reduced to 40. The training data was then fitted to a Random Forest model using 5-fold cross validation. The model was then used to the class variable in the test set and to answer the 20 quiz questions. The model correctly classified the 20 quiz questions.

Question

The data analysis in this report sets out to answer the following question.

Is it possible to classify the manner in which a dumbbell exercise was performed using sensor data from a glove, belt, arm-band and dumbbell?

Data

The training data for this project are available [here](#):

The test data are available [here](#):

The data for this project comes from [this source](#):

Citation: This report is based on data from the following paper:

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13) . Stuttgart, Germany: ACM SIGCHI, 2013.

Read more [here](#)

Assignment setup

The dataset has been placed in the directory `e:/rcode/chap08`

```
# Set the assignment working directory  
project_dir <- setwd("e:/rcode/chap08")
```

Libraries used in analysis

```
# load libraries  
library(dplyr, quietly = TRUE)  
library(caret, quietly = TRUE)  
library(YaleToolkit, quietly = TRUE)  
library(parallel, quietly = TRUE)  
library(doParallel)  
library(iterators)  
library(foreach)
```

Data Processing

This section describes (in words and code) how the data were loaded into R and processed for analysis. In particular, this shows how the analysis starts with a raw CSV file containing the data.

Preliminary data exploration

Read the first 5 rows and examine the data as it relates to the question

```
# Perform initial exploration  
  
# Peek at first 5 rows and find class of data  
peek5.train <- read.table("pml-training.csv", sep = ",", nrow = 5, header = TRUE)  
  
# examine variables  
str(peek5.train)  
  
# Peek at first 5 rows and find class of data  
peek5.test <- read.table("pml-testing.csv", sep = ",", nrow = 5, header = TRUE)  
  
# examine variables  
str(peek5.test)
```

Load all data

Load all data

```
# read data file  
har.all <- read.table("pml-training.csv", sep = ",", header = TRUE, na.strings = "NA", stringsAsFactors = FALSE)  
  
dim(har.all)
```

```
## [1] 19622 160
```

```
har.quiz <- read.table("pml-testing.csv", sep = ",", header = TRUE, na.strings = "NA", stringsAsFactors = FALSE)
dim(har.quiz)
```

```
## [1] 20 160
```

Use the `whatis()` function from the `YaleToolKit` to get additional perspective and insight on the data. Writing the data to a csv file and then browsing the data with Excel proved to be very useful.

```
var_explore_all <- whatis(har.all)

# create file var_exploration.csv
write.table(var_explore_all, file = "var_explore_all.csv", sep = ",", row.names = FALSE, col.names = TRUE)

var_explore_quiz <- whatis(har.quiz)

# create file var_exploration.csv
write.table(var_explore_quiz, file = "var_explore_quiz.csv", sep = ",", row.names = FALSE, col.names = TRUE)
```

Feature Extraction and Selection

1. Based on data exploration, there are 19622 observations across 160 variables in the `har.all`. However, reviewing the csv created from `var-exploration_all`, 67 of the variables all have 19216 missing values. For this reason, all of these variables will not be considered for the model
2. Based on **Section 5.1 Feature Extraction and Selection** from [the research paper](#), the following variables do not appear to have any roll in building the classification model: `new_window`, `raw_timestamp_part1`, `raw_timestamp_part2`, `cvtd_timestamp`, `user_name`, and `V1`. These variables will also not be considered for the model.
3. Additionally, Section 5.1 Feature Extraction and Selection makes no mention of using any of the kurtosis, or skewness variables. These variables will also not be considered for the model.

To further help with feature selection, for a given sensor, the following additional assumptions were made from reading Section 5.1 Feature Extraction and Selection:

Additional assumptions

- Gyro, implies pitch, roll and yaw in the variable names
- magnetometer, implies magnet in the variable names
- range, implies amplitude
- If gyro was mentioned as being used, all the raw variables for that sensor were included. For example, for the belt, it says “variance of the gyro” and therefore pitch, roll and yaw variables are included for that sensor
- If magnetometer was mentioned as being used, the x, y, and z magnet variables for that sensor were included in the model

Data Prep and Cleaning

From the data exploration, it can be seen that some variables need to have their data types coerced. For example, character to numeric and logical to numeric.

```
# convert classes

har.all$classe <- as.factor(har.all$classe)
har.all$max_yaw_belt <- as.numeric(har.all$max_yaw_belt)
har.quiz$avg_roll_belt <- as.numeric(har.quiz$avg_roll_belt)
har.quiz$var_roll_belt <- as.numeric(har.quiz$var_roll_belt)
har.quiz$var_total_accel_belt <- as.numeric(har.quiz$var_total_accel_belt)
har.quiz$max_roll_belt <- as.numeric(har.quiz$max_roll_belt)
har.quiz$min_roll_belt <- as.numeric(har.quiz$min_roll_belt)
har.quiz$max_pitch_belt <- as.numeric(har.quiz$max_pitch_belt)
har.quiz$amplitude_pitch_belt <- as.numeric(har.quiz$amplitude_pitch_belt)
har.quiz$var_accel_arm <- as.numeric(har.quiz$var_accel_arm)
har.quiz$var_pitch_dumbbell <- as.numeric(har.quiz$var_pitch_dumbbell)
har.quiz$var_roll_dumbbell <- as.numeric(har.quiz$var_roll_dumbbell)
har.quiz$var_yaw_dumbbell <- as.numeric(har.quiz$var_yaw_dumbbell)
har.quiz$amplitude_roll_dumbbell <- as.numeric(har.quiz$amplitude_roll_dumbbell)

#dumbbell
har.quiz$min_pitch_dumbbell <- as.numeric(har.quiz$min_pitch_dumbbell)
har.quiz$min_roll_dumbbell <- as.numeric(har.quiz$min_roll_dumbbell)
har.quiz$min_yaw_dumbbell <- as.numeric(har.quiz$min_yaw_dumbbell)
har.quiz$max_pitch_dumbbell <- as.numeric(har.quiz$max_pitch_dumbbell)
har.quiz$max_roll_dumbbell <- as.numeric(har.quiz$max_roll_dumbbell)
har.quiz$max_yaw_dumbbell <- as.numeric(har.quiz$max_yaw_dumbbell)

#arm
har.quiz$min_pitch_arm <- as.numeric(har.quiz$min_pitch_arm)
har.quiz$min_roll_arm <- as.numeric(har.quiz$min_roll_arm)
har.quiz$min_yaw_arm <- as.numeric(har.quiz$min_yaw_arm)
har.quiz$max_pitch_arm <- as.numeric(har.quiz$max_pitch_arm)
har.quiz$max_roll_arm <- as.numeric(har.quiz$max_roll_arm)
har.quiz$max_yaw_arm <- as.numeric(har.quiz$max_yaw_arm)

#belt
har.quiz$min_pitch_belt <- as.numeric(har.quiz$min_pitch_belt)
har.quiz$min_roll_belt <- as.numeric(har.quiz$min_roll_belt)
har.quiz$min_yaw_belt <- as.numeric(har.quiz$min_yaw_belt)
har.quiz$max_pitch_belt <- as.numeric(har.quiz$max_pitch_belt)
har.quiz$max_roll_belt <- as.numeric(har.quiz$max_roll_belt)
har.quiz$max_yaw_belt <- as.numeric(har.quiz$max_yaw_belt)

#forearm
har.quiz$min_pitch_forearm <- as.numeric(har.quiz$min_pitch_forearm)
har.quiz$min_roll_forearm <- as.numeric(har.quiz$min_roll_forearm)
har.quiz$min_yaw_forearm <- as.numeric(har.quiz$min_yaw_forearm)
har.quiz$max_pitch_forearm <- as.numeric(har.quiz$max_pitch_forearm)
har.quiz$max_roll_forearm <- as.numeric(har.quiz$max_roll_forearm)
har.quiz$max_yaw_forearm <- as.numeric(har.quiz$max_yaw_forearm)
```

```

# amplitude train
har.all$amplitude_yaw_arm <- as.numeric(har.all$amplitude_yaw_arm)
har.all$amplitude_yaw_belt <- as.numeric(har.all$amplitude_yaw_belt)
har.all$amplitude_yaw_forearm <- as.numeric(har.all$amplitude_yaw_forearm)

# amplitude test
har.quiz$amplitude_yaw_arm <- as.numeric(har.quiz$amplitude_yaw_arm)
har.quiz$amplitude_yaw_belt <- as.numeric(har.quiz$amplitude_yaw_belt)
har.quiz$amplitude_yaw_forearm <- as.numeric(har.quiz$amplitude_yaw_forearm)

# train yaw
har.all$max_yaw_belt <- as.numeric(har.all$max_yaw_belt)
har.all$max_yaw_dumbbell <- as.numeric(har.all$max_yaw_dumbbell)
har.all$max_yaw_forearm <- as.numeric(har.all$max_yaw_forearm)
har.all$min_yaw_belt <- as.numeric(har.all$min_yaw_belt)
har.all$min_yaw_dumbbell <- as.numeric(har.all$min_yaw_dumbbell)
har.all$min_yaw_forearm <- as.numeric(har.all$min_yaw_forearm)

```

Creating the training and test datasets

Even though this assignment implies a specific training and test dataset, this analysis report partitions the data differently. For this report, the original test dataset is being used solely for the answering the quiz questions and is called `har.quiz`.

The original training dataset has been renamed to `har.all`. `har.all` is further divided into a new train and test dataset, `har.train` and `har.test`, respectively.

```

# create a training and test data set from the har.all data
set.seed(3465)
inTrain <- createDataPartition(y = har.all$classe, p=0.75, list=FALSE)

har.train <- har.all[inTrain,]
har.test <- har.all[-inTrain,]

```

Model Computation

```

# create a vector of variables bases on defined feature selection criteria
har.train.tmp <- select(har.train,
  classe,
  roll_belt,
  pitch_belt,
  yaw_belt,
  total_accel_belt,
  gyros_belt_x,
  gyros_belt_y,
  gyros_belt_z,
  accel_belt_x,
  accel_belt_y,
  accel_belt_z,
  magnet_belt_x,
  magnet_belt_y,

```

```

magnet_belt_z,
gyros_forearm_x,
gyros_forearm_y,
gyros_forearm_z,
accel_forearm_x,
accel_forearm_y,
accel_forearm_z,
magnet_forearm_x,
magnet_forearm_y,
magnet_forearm_z,
gyros_arm_x,
gyros_arm_y,
gyros_arm_z,
accel_arm_x,
accel_arm_y,
accel_arm_z,
magnet_arm_x,
magnet_arm_y,
magnet_arm_z,
gyros_dumbbell_x,
gyros_dumbbell_y,
gyros_dumbbell_z,
accel_dumbbell_x,
accel_dumbbell_y,
accel_dumbbell_z,
magnet_dumbbell_x,
magnet_dumbbell_y,
magnet_dumbbell_z
)

```

Citation:

Information on how to improve performance of Random Forest in caret came from the following: https://github.com/josh1elbaum/parallel_rf

setup and register cluster

```

cluster <- makeCluster(detectCores() - 1) # convention to leave 1 core for OS
registerDoParallel(cluster)

```

configure train control parameters to include cross validation parallel computation

```

fitControl <- trainControl(method = "cv",
                           number = 5,
                           allowParallel = TRUE)

```

train the model

```

modFit2 <- train(classe ~ ., method="rf", trControl = fitControl, data=har.train.tmp)

```

shutdown the cluster

```

stopCluster(cluster)

```

Results and Analysis

Show model results and use the model to predict the classe variable on the test data.

```
# Show model results
modFit2
```

```
## Random Forest
##
## 14718 samples
## 40 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 11775, 11775, 11774, 11774, 11774
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.9902161 0.9876227 0.0016005094 0.0020258318
## 21 0.9890610 0.9861608 0.0019868724 0.0025144113
## 40 0.9858676 0.9821223 0.0007829572 0.0009904767
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
# use model to predict the classe variable on the test data
pred.modFit2.test <- predict(modFit2, newdata = har.test)

# generate confusion matrix and estimate of out-of-sample error
conf.matrix.test <- confusionMatrix(data = pred.modFit2.test, har.test$classe)
conf.matrix.test
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A  B  C  D  E
## A 1395    1  0  1  0
## B   0  946  6  0  0
## C   0   2 848 11  0
## D   0   0  1 791  0
## E   0   0  0   1 901
##
## Overall Statistics
##
##           Accuracy : 0.9953
##           95% CI : (0.993, 0.997)
##       No Information Rate : 0.2845
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9941
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity    1.0000  0.9968  0.9918  0.9838  1.0000
```

## Specificity	0.9994	0.9985	0.9968	0.9998	0.9998
## Pos Pred Value	0.9986	0.9937	0.9849	0.9987	0.9989
## Neg Pred Value	1.0000	0.9992	0.9983	0.9968	1.0000
## Prevalence	0.2845	0.1935	0.1743	0.1639	0.1837
## Detection Rate	0.2845	0.1929	0.1729	0.1613	0.1837
## Detection Prevalence	0.2849	0.1941	0.1756	0.1615	0.1839
## Balanced Accuracy	0.9997	0.9977	0.9943	0.9918	0.9999

Model has an overall accuracy of 0.99531

What are the most important variables in the model?

```
modFit2.importance <- as.data.frame(modFit2$finalModel$importance)

# move row names to a new column
modFit2.importance$new <- rownames(modFit2.importance)
rownames(modFit2.importance) <- NULL

modFit2.importance <- arrange(modFit2.importance, desc(MeanDecreaseGini))
modFit2.importance
```

##	MeanDecreaseGini	new
## 1	633.5613	roll_belt
## 2	562.0207	yaw_belt
## 3	543.2001	magnet_dumbbell_z
## 4	493.4976	magnet_dumbbell_y
## 5	474.0261	pitch_belt
## 6	440.2459	magnet_dumbbell_x
## 7	390.3627	accel_dumbbell_y
## 8	369.0083	magnet_belt_z
## 9	357.0062	accel_belt_z
## 10	353.0827	accel_dumbbell_z
## 11	339.1845	magnet_belt_y
## 12	320.1377	accel_dumbbell_x
## 13	313.3496	accel_forearm_x
## 14	312.2319	magnet_arm_x
## 15	305.9672	accel_arm_x
## 16	292.3523	magnet_arm_y
## 17	290.1414	magnet_forearm_x
## 18	283.4474	gyros_belt_z
## 19	272.6965	magnet_forearm_z
## 20	271.2225	gyros_dumbbell_y
## 21	268.3218	magnet_forearm_y
## 22	260.1735	accel_forearm_z
## 23	252.9213	magnet_arm_z
## 24	252.2037	magnet_belt_x
## 25	230.0715	accel_arm_y
## 26	226.1951	gyros_arm_x
## 27	224.4454	accel_forearm_y
## 28	222.9943	total_accel_belt
## 29	215.9327	accel_arm_z
## 30	210.8938	gyros_arm_y


```
## 31      208.5369   gyros_forearm_y
## 32      191.7140 gyros_dumbbell_x
## 33      184.2749   accel_belt_x
## 34      169.5078   accel_belt_y
## 35      165.8144   gyros_belt_x
## 36      156.0640 gyros_forearm_z
## 37      155.7811 gyros_forearm_x
## 38      152.5882   gyros_belt_y
## 39      142.7339 gyros_dumbbell_z
## 40      126.7566   gyros_arm_z
```

Use model to predict quiz answers

```
pred.modFit2.quiz <- predict(modFit2, newdata = har.quiz)
pred.modFit2.quiz
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Conclusion

This report showed how a dataset without a code book could be explored and analyzed to build a model to classify how an activity was performed. The model used 5-fold cross validation and showed out-of-sample accuracy/error on the test data. Also, the model correctly classified the 20 quiz questions.

However, this model would not generalize well as (1) the HAR experiment was carried out in a very controlled environment. (2) There were not enough participants in the study.